

STAT 6358 Project Proposal - Adopting the CAMP Framework of Censoring Zeros for RNA-Seq.

Callum Doyle

2025-03-30

Abstract

RNA-seq data is sparse and an ongoing challenge in this field is the choice of zero-handling method when the goal is to test for differential expression. CAMP (Chan and Li (2024)) is a recently developed method that treats zero counts as censored values, facilitating the use of survival analysis methods for differential abundance analysis (DAA). The goal of this project is to assess the feasibility of treating zero counts as censored values in RNA-seq data, in contrast to traditional methods which remove lowly-expressed genes or add a pseudocount to zeros. In particular, we aim to investigate the impact on differential analysis of RNA-seq data when zeros are treated as censored by exploring the following interesting questions. 1. Does treating zeros as censored change the differentially expressed genes (DEGs) results meaningfully? 2. Do the unique DEGs detected by the CAMP framework make sense biologically? 3. Under what conditions (for example, sample size or sparsity) is there a large overlap between CAMP and more-traditional RNA-seq methods, and equivalently, when do they disagree? By answering these questions, we aim to determine if treating zeros as censored is a valid approach that could, in-turn, open up new ideas to existing methodology by including the idea that a zero-count is partial information. Although CAMP was developed for microbiome data, the authors suggest that it could extend to RNA-seq data due to the similarity in sparsity.

This project will be uploaded to <https://github.com/cjd04/stat6358project>. Ideally, it will include the simulation and real data analysis code, the simulated datasets, references, and whatever else that might be relevant (if space/memory limits allow, particularly for the datasets).

Proposal

Introduction

The nature of high-throughput sequencing data leads to high sparsity. Many zero counts in a dataset are problematic and many methods have been proposed to tackle the challenges that sparse datasets pose. There is still no consensus for the best choice of zero-handling method, and the choice appeared to be based on preference and little else. Silverman et al. (2020) explored different zero-handling methods across a variety of scenarios (simulated and real datasets) and demonstrated that the choice of method drastically influences the outcome. Additionally, the behaviour of the different methods was examined under different zero-generating processes (ZGPs). ZGPs were categorised into the following four categories. The first category is sampling zeros, this is when a zero count for a given sample is not due to absence, but due to a sampling effect. The second category is biological zeros, this is when a zero count is truly zero. We consider that a sequence with a zero count is absent from the biological system. The third and fourth categories are technical zeros, partial and complete; where the technical bias either partially or completely inhibits measurement. Although the purpose of this project is not to investigate how well methods handle different ZGPs, it is useful to keep the types of zeros in mind when examining results for meaningful interpretation. Overall, there is still a gap in the literature for the best way of handling zeros. It is reasonable to assume that not all zeros are the same. This gives way to the idea that some zeros actually contain partial information.

The framework introduced by Chan and Li (2024), Censoring-based analysis of Microbiome Proportions (CAMP), introduces the idea of treating zero counts as censored observations, then transforming the count matrix into data that can be analysed by survival analysis methods. The idea is to treat the zero counts as partially observed data and perform some transformations to convert the data to time-to-event-like data that is likely to be without ties, which addresses the low-power issue that nonparametric tests face in the presence of ties. CAMP assumes compositionality of the data and results show that type 1 is well-controlled and powerful. The authors suggest the CAMP framework can be applied to RNA-seq data, if treated as compositional. We propose to adapt the CAMP framework to RNA-seq data and compare the results to three methods developed for RNA-seq, edgeR, DESeq2, and ALDEx2. edgeR and DESeq2 are more traditional methods that do not assume relative abundance, whereas ALDEx2 assumes the data is compositional

CAMP approach

The general idea of CAMP is as follows. Once the raw count matrix $X = \{x_{ij}\}$ has been obtained (OTU table in the case of CAMP), where $i = 1, \dots, n$ indexes the samples and $j = 1, \dots, p$ indexes the taxa (translation for RNA-seq data: genes), define a surrogate read count matrix and an indicator matrix $X^* = \{x_{ij}^*\}$ and $\Delta = \{\delta_{ij}\}$ such that $x_{ij}^* = x_{ij}$ if $x_{ij} > d$

or $x_{ij}^* = d$ if $x_{ij} = 0$, and $\delta_{ij} = 1$ if $x_{ij} > d$ and 0 otherwise, where $d > 0$ is a predefined detection limit in a sequencing study. Typically, d is just the minimum non-zero count across the whole count matrix. The next step is to get the relative abundance matrix which is formed from the surrogate matrix, so it is left-censored and there are no zero counts. Next, the relative abundance matrix is transformed into right-censored time-to-event data by applying a negative log-transformation. This new “time” matrix facilitates the use of survival analysis techniques and the analogous interpretation (of whether the event has occurred or not) for differential gene expression is based on the presence or absence of a taxon (gene) at a specific relative abundance level. In survival analysis, we are interested in the “at-risk” population at a given time point which consists of the subjects (taxa) who have not yet experienced the event. For us, this translates to: at a given cut-off point of relative abundance, the taxa (gene) is either present (differentially expressed) or absent (not DE). This creates a 2x2 table of presence/absence for condition 1 vs 2, and a log-rank test can be conducted to test for significance. Note that no distributional assumptions are required for this test, and the transformed data is likely to be without ties, which Chan and Li (2024) acknowledge tackles the issue of ties causing reduced power in nonparametric tests.

Preliminary results

So far, we have conducted simulations under various scenarios comparing the four methods. Data generation was conducted using `compcoder` (Soneson (2014)), where options for generation are straightforward, which makes it ideal for conducting multiple simulations. We assumed the default setting for each method; the reasoning for this is that the most optimal setting for each method under a particular scenario is not of great interest. Indeed, the interest lies in the capability of the CAMP framework to outperform other methods that were originally designed for RNA-seq data. That is, if CAMP performs similarly or better than the other traditional methods, it is reasonable to believe that treating zero counts as censored is a valid approach that would be worth exploring in greater detail.

We conducted a simulation study to compare the performance of CAMP with the three traditional methods. All settings assumed two conditions (group 1 and group 2), where the following settings were modified to produce a total of 18 simulation scenarios.

- The initial number of genes in dataset: 10,000; 15,000; 20,000.
- The number of samples in each condition, $\frac{n}{2}$: 10; 25; 50.
- The proportion of truly differentially expressed genes: 0.05; 0.1.

The simulation results suggest that CAMP performs similarly to the other RNA-seq methods. Indeed, Figure 1 demonstrates that the power of CAMP follows the same pattern as the other methods across simulation settings, and although CAMP was never the most powerful method, it was consistently the second or third-best performing, and with a large enough sample size was often the second-best. Additionally, CAMP had a well-controlled type 1 error rate (see

Table 1 in the appendix). Performance appears to primarily depend on sample size and the number of true DEGs, rather than the initial number of genes in the dataset.

These results are promising and suggest that further investigations should be conducted. In particular, extending the simulation to include different levels of sparsity –perhaps by manually inflating the zero count from the generated RNA-seq data– will be an important benchmark for comparison. Moreover, a second data-generation method could be useful to protect against any underlying biases that the data-generating mechanism may have. Finally, a real data analysis is important to test the biological plausibility of the results produced by treating zeros as censored. It would be good to explore two types of real RNA-seq data, a whole/bulk RNA-seq data set, and a single-cell RNA-seq dataset. Choosing a dataset that has been well-explored and biologically interpretable is ideal, as it will enable us to investigate if the resulting DEGs found by a censored-zero approach are meaningful.




Figure 1: Power table displaying results across different simulation scenarios. There are two numbers above each chart: the top number is the total number of initial genes in the dataset; the bottom number is the proportion of true DEGs. For example, there are 1000 DEGs in the second chart in the top row.

Appendix

Table 1: Power and type 1 error table across each simulation scenario. The output is raw because the plot is much easier to understand and the type 1 error is pretty good all across the board for these settings (except for edgeR later on).

method	power	Type1 error	n.genes	n.per.grp	n.trueDEGs
CAMP	0.522	0.004	10000	10	0.05
edgeR	0.538	0.031	10000	10	0.05
DESeq2	0.608	0.005	10000	10	0.05
ALDEx2	0.438	0.001	10000	10	0.05
CAMP	0.472	0.003	15000	10	0.05
edgeR	0.528	0.030	15000	10	0.05
DESeq2	0.580	0.004	15000	10	0.05
ALDEx2	0.416	0.000	15000	10	0.05
CAMP	0.488	0.004	20000	10	0.05
edgeR	0.523	0.030	20000	10	0.05
DESeq2	0.597	0.005	20000	10	0.05
ALDEx2	0.411	0.000	20000	10	0.05
CAMP	0.708	0.004	10000	25	0.05
edgeR	0.730	0.041	10000	25	0.05
DESeq2	0.796	0.003	10000	25	0.05
ALDEx2	0.674	0.001	10000	25	0.05
CAMP	0.676	0.004	15000	25	0.05
edgeR	0.671	0.038	15000	25	0.05
DESeq2	0.787	0.003	15000	25	0.05
ALDEx2	0.620	0.001	15000	25	0.05
CAMP	0.676	0.003	20000	25	0.05
edgeR	0.698	0.039	20000	25	0.05
DESeq2	0.783	0.003	20000	25	0.05
ALDEx2	0.630	0.001	20000	25	0.05
CAMP	0.816	0.005	10000	50	0.05
edgeR	0.804	0.047	10000	50	0.05
DESeq2	0.904	0.004	10000	50	0.05
ALDEx2	0.802	0.002	10000	50	0.05
CAMP	0.801	0.013	15000	50	0.05
edgeR	0.779	0.044	15000	50	0.05
DESeq2	0.885	0.003	15000	50	0.05
ALDEx2	0.757	0.001	15000	50	0.05
CAMP	0.813	0.003	20000	50	0.05
edgeR	0.760	0.043	20000	50	0.05
DESeq2	0.895	0.003	20000	50	0.05

method	power	Type1 error	n.genes	n.per.grp	n.trueDEGs
ALDEx2	0.763	0.001	20000	50	0.05
CAMP	0.522	0.008	10000	10	0.10
edgeR	0.579	0.068	10000	10	0.10
DESeq2	0.634	0.007	10000	10	0.10
ALDEx2	0.462	0.001	10000	10	0.10
CAMP	0.509	0.006	15000	10	0.10
edgeR	0.543	0.064	15000	10	0.10
DESeq2	0.601	0.006	15000	10	0.10
ALDEx2	0.433	0.001	15000	10	0.10
CAMP	0.495	0.006	20000	10	0.10
edgeR	0.524	0.062	20000	10	0.10
DESeq2	0.588	0.007	20000	10	0.10
ALDEx2	0.417	0.001	20000	10	0.10
CAMP	0.736	0.007	10000	25	0.10
edgeR	0.751	0.089	10000	25	0.10
DESeq2	0.821	0.006	10000	25	0.10
ALDEx2	0.712	0.002	10000	25	0.10
CAMP	0.713	0.007	15000	25	0.10
edgeR	0.723	0.085	15000	25	0.10
DESeq2	0.810	0.006	15000	25	0.10
ALDEx2	0.679	0.001	15000	25	0.10
CAMP	0.697	0.007	20000	25	0.10
edgeR	0.686	0.080	20000	25	0.10
DESeq2	0.779	0.005	20000	25	0.10
ALDEx2	0.656	0.001	20000	25	0.10
CAMP	0.823	0.009	10000	50	0.10
edgeR	0.815	0.097	10000	50	0.10
DESeq2	0.904	0.007	10000	50	0.10
ALDEx2	0.802	0.002	10000	50	0.10
CAMP	0.821	0.013	15000	50	0.10
edgeR	0.782	0.092	15000	50	0.10
DESeq2	0.905	0.005	15000	50	0.10
ALDEx2	0.797	0.001	15000	50	0.10
CAMP	0.845	0.009	20000	50	0.10
edgeR	0.775	0.091	20000	50	0.10
DESeq2	0.908	0.005	20000	50	0.10
ALDEx2	0.796	0.002	20000	50	0.10

References

- Chan, Lap Sum, and Gen Li. 2024. “Zero Is Not Absence: Censoring-Based Differential Abundance Analysis for Microbiome Data.” *Bioinformatics* 40 (2): btae071.
- Silverman, Justin D, Kimberly Roche, Sayan Mukherjee, and Lawrence A David. 2020. “Naught All Zeros in Sequence Count Data Are the Same.” *Computational and Structural Biotechnology Journal* 18: 2789–98.
- Soneson, Charlotte. 2014. “compcodeR?an r Package for Benchmarking Differential Expression Methods for RNA-Seq Data.” *Bioinformatics* 30 (17): 2517–18.